



Introduction

25 February 2018

Henry R.P

- **Mapr Certified Hadoop Administrator**
- **IBM Certified Application Developer**
- **IBM Certified Solution Designer**
- **SAP Certified Development Consultant.**

IT Architect , Author & Corporate trainer
15 +Year of IT Experience

- **TOGAF – Enterprise Architect**
- **CIPM – Certificate in Project Management.**

Training & Consulting on:

NOSQL & Bigdata – Hadoop, Couchbase, Cassandra, MongoDB, CDH, BigInsight

Predictive Analytics – R & SAS

EAI:- Mule / Fuse ESB / Spring Integration / JBI / Apache Camel / Talend / Apache Service Mix


Portal:- Liferay, SAP Netweaver.

Application server:- WAS, Tomcat, WebLogic, Jboss

Architecture: EA, TOGAF, CoBIT etc.

JEE Framework

OSGI - Eclipse PDE/Equinox/Virgo/Spring DM/ Felix / Karaf



Clientele



THOMSON REUTERS



GE Healthcare



Introduce Yourself.

Name

Year of Experience.

Skills Level

RDMS

NoSql /Couchbase

Expectation, if any.

Outline

Overview - NoSQL Basic
Couchbase Server Architecture
Couchbase Administration - Webconsole
Bucket
Document Database Basic
Upgrade
Couchbase SDK - Overview
Cluster Administration
Views
N1QL
Security - LDAP
Client API - Java API , Rx and Spring DB
XDCR
FTS
Monitoring + Back up.

Time	
9.30 – 11.00 AM	Session I
11.00 AM to 11.15 AM	Tea Break
11.15 AM to 12.45 PM	Session II
12.45 PM to 1.45 PM	Lunch Break
1.45 PM to 3.15 PM	Session III
3.15 PM to 3.30 PM	Tea Break
3.30 PM to 5.30 PM	Session IV



An introduction to **NoSQL** databases

Relational databases

5

Benefits of Relational databases:

- Designed for all purposes
- ACID
- Strong consistency, concurrency, recovery
- Mathematical background
- Standard Query language (SQL)
- Lots of tools to use with i.e: Reporting services, entity frameworks, ...
- Vertical scaling (upscaling)

Object / Object-relational databases were not practical. Mainly because of Impedance mismatch

Era of Distributed Computing


6

But...

- ❑ Relational databases were not built for **distributed applications**.

Because...

- ❑ Joins are expensive
- ❑ Hard to scale horizontally
- ❑ Impedance mismatch occurs
- ❑ Expensive (product cost, hardware, Maintenance)



Era of Distributed Computing

7

But...

- ❑ Relational databases were not built for **distributed applications**.


Because...

- ❑ Joins are expensive
- ❑ **Hard to scale horizontally**
- ❑ Impedance mismatch occurs
- ❑ Expensive (product cost, hardware, Maintenance)

And....

It's weak in:

- ❑ Speed (performance)
- ❑ High availability
- ❑ Partition tolerance



Characteristics of NoSQLdatabases

12

- ❑ Non relational
- ❑ Cluster friendly
- ❑ Schema-less
- ❑ 21 century web
- ❑ Open-source



Characteristics of NoSQL databases

13

NoSQL avoids:

- ❑ Overhead of ACID transactions
- ❑ Complexity of SQL query
- ❑ Burden of up-front schema design
- ❑ DBA presence
- ❑ Transactions (It should be handled at application layer)

Provides:

- ❑ Easy and frequent changes to DB
- ❑ Horizontal scaling (scaling out)
- ❑ Solution to Impedance mismatch
- ❑ Fast development



What is a schema-less data model?

15

```
create table customers (id int, firstname text, lastname text)
insert into customers (firstname, middlename, lastname) values (...)
```

In relational Databases:

- ❑ You can't add a record which does not fit the schema
- ❑ You need to add NULLs to unused items in a row
- ❑ We should consider the datatypes. i.e : you can't add a string to an integer field
- ❑ You can't add multiple items in a field (You should create another table: primary-key, foreign key, joins, normalization, ... !!!)




What is a schema-less data model?

16

In NoSQL Databases:


- ❑ There is no schema to consider
- ❑ There is no unused cell
- ❑ There is no datatype (implicit)
- ❑ Most of considerations are done in application layer
- ❑ We gather all items in an aggregate (document)

Relational Model



Document Model

Collection ("Things")



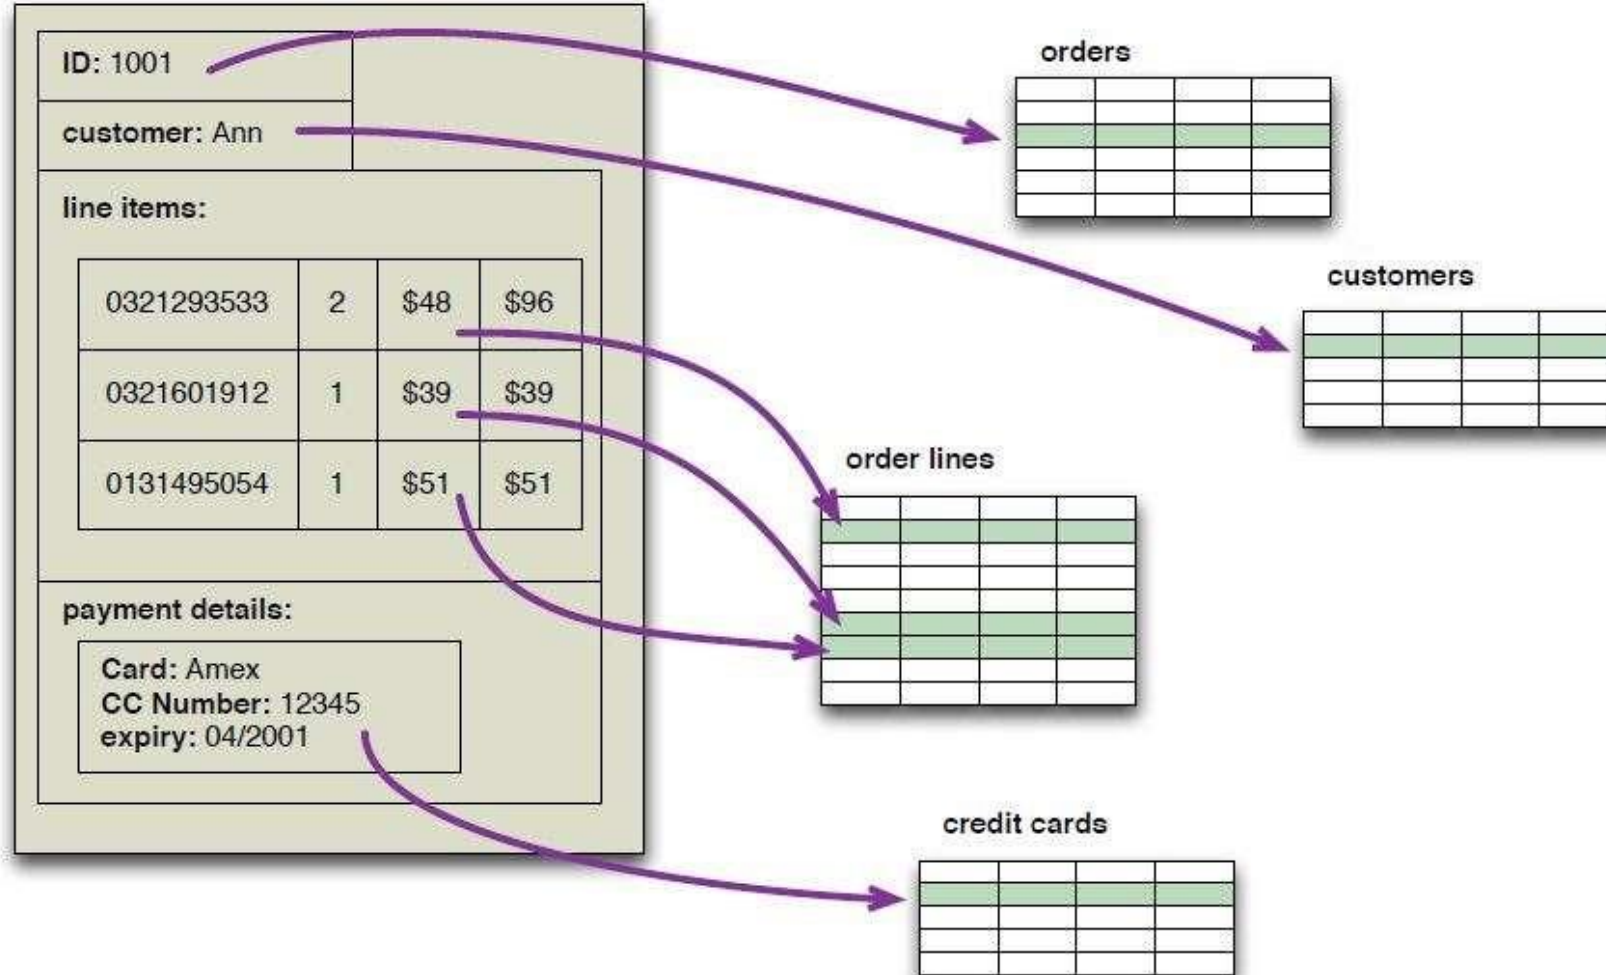
What is Aggregation?

17

- ❑ The term comes from Domain Driven Design
- ❑ Shared nothing architecture
- ❑ An aggregate is a cluster of domain objects that can be treated as a single unit
- ❑ Aggregates are the basic element of transfer of data storage - you request to load or save whole aggregates
- ❑ Transactions should not cross aggregate boundaries
- ❑ This mechanism reduces the join operations to a minimal level

What is Aggregation?

18



What is Aggregation?

19

```
{
  "id": "1001",
  "firstName": "Ann",
  "lastName": "Williams",
  "age": 55,
  "purchasedItems":
  {
    0321290533 {qty, price...}
    0321601912 {qty, price...}
    0131495054 {qty, price...}
  }
  "paymentDetails":
  { cc info... }
  "address":
  {
    "street": "1234 Park",
    "city": "San Francisco",
    "state": "CA",
    "zip": "94102"
  }
}
```

ID: 1001			
customer: Ann			
line items:			
0321293533	2	\$48	\$96
0321601912	1	\$39	\$39
0131495054	1	\$51	\$51
payment details:			
Card: Amex CC Number: 12345 expiry: 04/2001			

aggregate



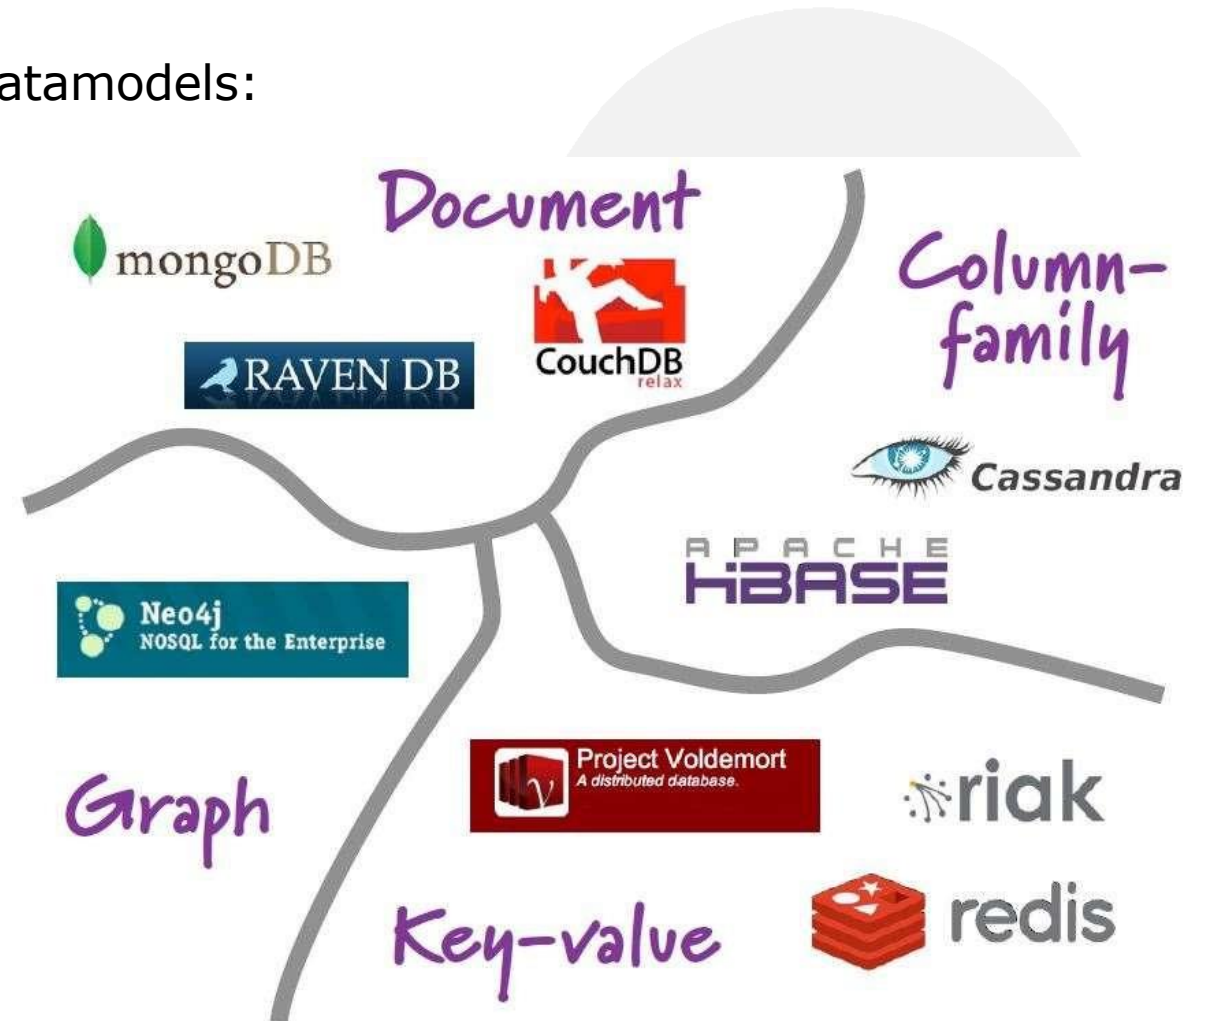
Aggregate Data Models

21

NoSQL databases are classified in four major datamodels:

- Key-value
- Document
- Column family
- Graph

Each DB has its own query language



Key-value data model

22

- The main idea is the use of a hash table
- Access data (values) by strings called keys
- Data has no required format – data may have any format
- Data model: (key, value) pairs
- Basic Operations:
Insert(key,value), Fetch(key), Update(key), Delete(key)



Car	
Key	Attributes
1	Make: Nissan Model: Pathfinder Color: Green Year: 2003
2	Make: Nissan Model: Pathfinder Color: Blue Color: Green Year: 2005 Transmission: Auto

Column family data model

24

- The column is lowest/smallest instance of data.
It is a tuple that contains a name, a value and a timestamp



ColumnFamily: Authors		
Key	Value	
"Eric Long"	Columns	
	Name	Value
	"email"	"eric (at) long.com"
	"country"	"United Kingdom"
	"registeredSince"	"01/01/2002"
"John Steward"	Columns	
	Name	Value
	"email"	"john.steward (at) somedomain.com"
	"country"	"Australia"
	"registeredSince"	"01/01/2009"
"Ronald Mathies"	Columns	
	Name	Value
	"email"	"ronald (at) sodeso.nl"
	"country"	"Netherlands, The"
	"registeredSince"	"01/01/2010"

Graph data model

26

- ❑ Based on Graph Theory.
- ❑ Scale vertically, no clustering.
- ❑ You can use graph algorithms easily
- ❑ Transactions
- ❑ ACID



Document-based datamodel

27

- Usually JSON like interchange model.
- Query Model: JavaScript-like or custom.
- Aggregations: **Map/Reduce**
- Indexes are done via B-Trees.
- unlike simple key-value stores, both keys and values are fully searchable in document databases.

```
{  
  person: {  
    first_name: "Peter",  
    last_name: "Peterson",  
    addresses: [  
      {street: "123 Peter St"},  
      {street: "504 Not Peter St"}  
    ],  
  },  
}
```



mongoDB

What we need?

34

- We need a distributed database system having such features:
 - **–Fault tolerance**
 - **–High availability**
 - **–Consistency**
 - **–Scalability**

What we need?

35

- We need a distributed database system having such features:
 - **–Fault tolerance**
 - **–High availability**
 - **–Consistency**
 - **–Scalability**

Which is impossible!!!
According to CAP theorem

Polyglot persistence : the future of database systems

45

The future is:

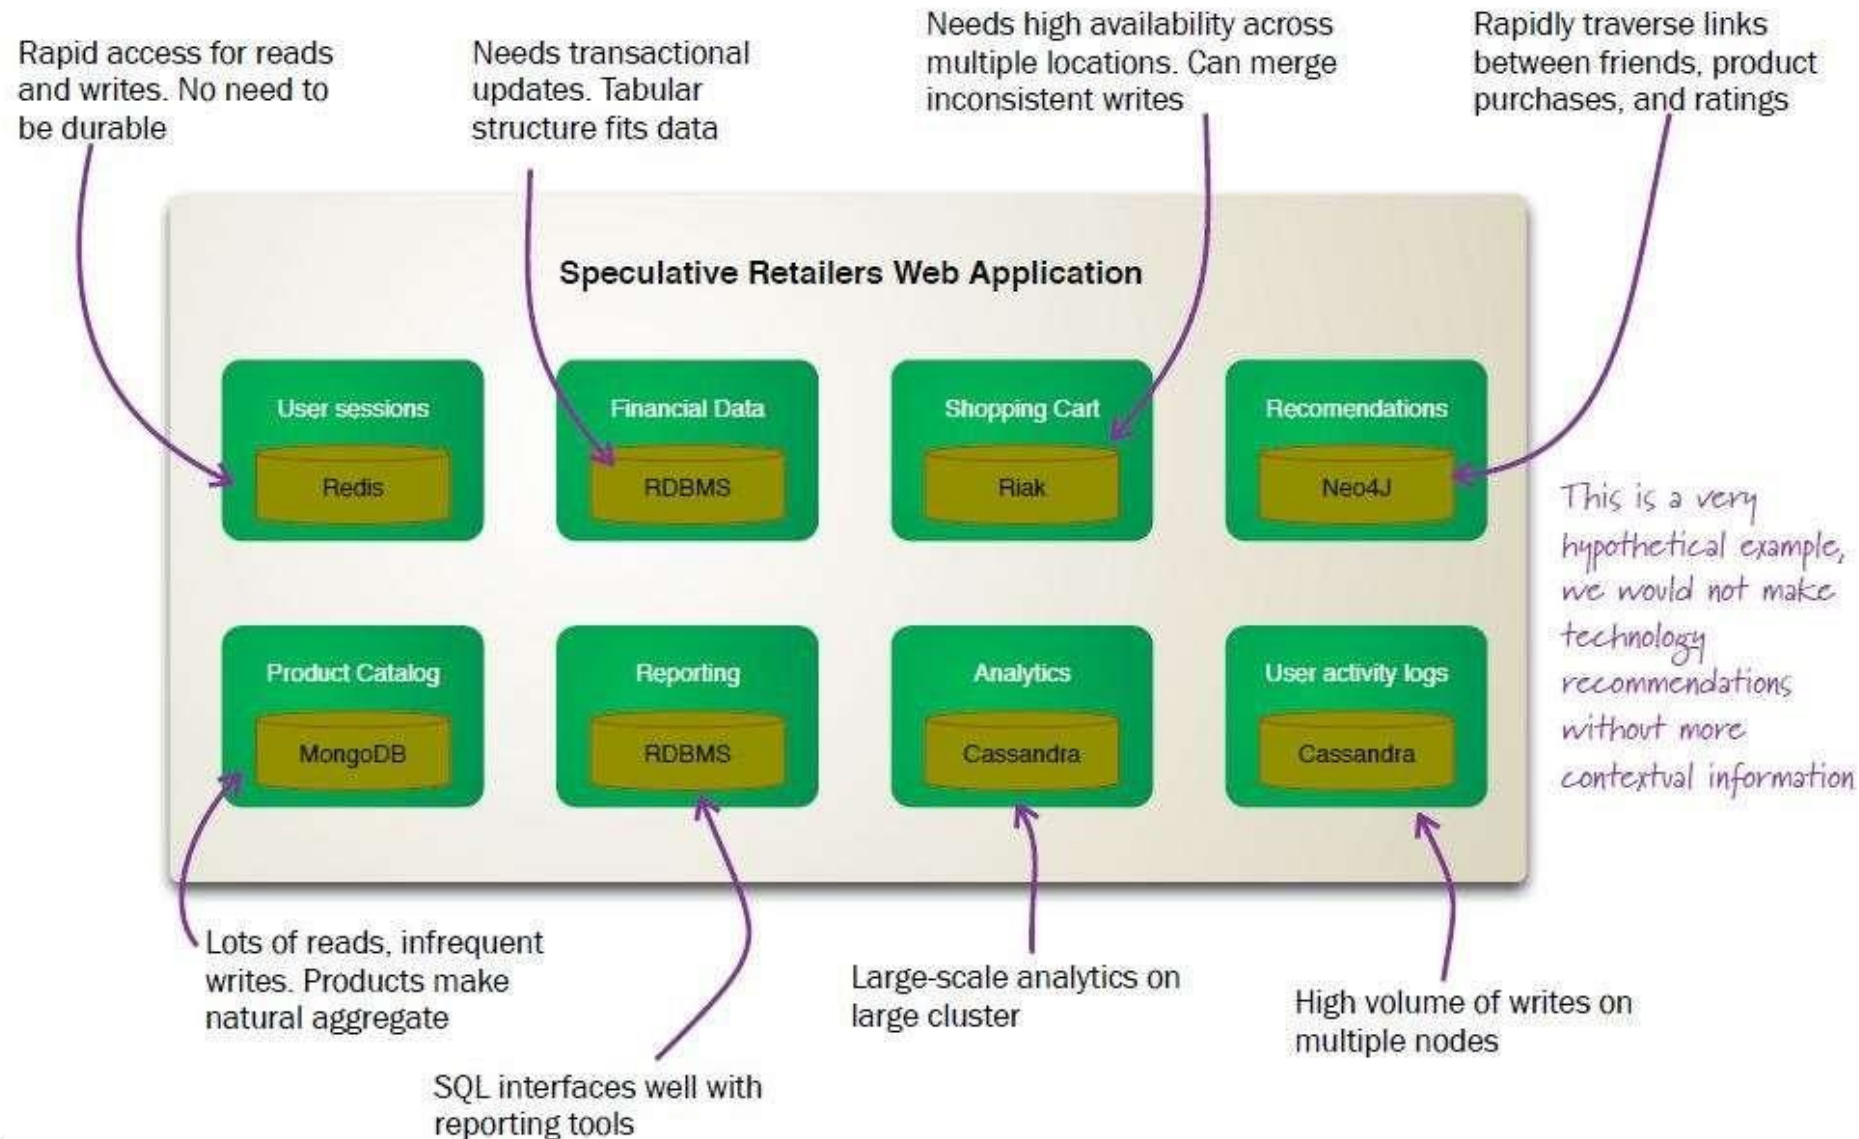
~~NoSQL Databases~~

Polyglot Persistence

- Future databases are the combination of SQL & NoSQL
- We still need relational databases

Overview of a polygot db

46



Conclusion:

49

Before you choose NoSQL as a solution:

